

```

/**
 * Matrix Multiplication.
 * Consider matrices A (L,M)    (L lines, M cols)
 *                        and B (M,N)    (M lines, N cols)
 * we want to compute resulting matrix
 *                        A x B = C (L,N)
 *
 * The input should be in a text file formatted as:
 * A : [0 1 2 3 4] [5 6 7 8 9]
 * B : [0 1 2] [3 4 5] [6 7 8] [9 10 11] [12 13 14]
 *
 * Stéphane Genaud, Nov 2013
 */

import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;

public class Mm {

    public static class Map extends Mapper<LongWritable, Text, Text, Text> {

        /**
         * map : the map function implementation
         *
         * @param key
         * @param value the data from input files, one line per map
         * invocation
         * @param context a place to write the (key,value) pairs we want to
         * emit
         */
        public void map(LongWritable key, Text value, Context context)
            throws IOException, InterruptedException {

            Text outputKey = new Text();
            Text outputValue = new Text();
            Configuration conf = context.getConfiguration();
            int L = Integer.parseInt(conf.get("L"));
            int M = Integer.parseInt(conf.get("M"));
            int N = Integer.parseInt(conf.get("N"));
            String [][] A = new String [L][M];
            String [][] B = new String [M][N];
            // separate matrix anem and values
            String [] matStruct = value.toString().split(":");
            matStruct[0]= matStruct[0].trim();
            matStruct[1]= matStruct[1].trim();
            //split values in tokens, each token being a matrix line
            String [] matLines = matStruct[1].split("\\\\");

            // A
            if (matStruct[0].equals("A")) {
                // ----- first parse matrix values and store in an array
                for (int i=0;i<L;i++) {

```

```

                    // remove leading '|'
                    // (replace doesn't use a regexp so no protection)
                    matLines[i] = matLines[i].replace("[", "").trim();
                    System.out.println(matLines[i]);
                    A[i] = matLines[i].split(" ");
                }
            }
            // ----- second, emit the values needed for products
            // emit (key,value)
            for (int i=0;i<L;i++) {
                for (int j=0;j<N;j++) {
                    for (int k=0;k<M;k++) {
                        System.out.println("C_"+i+"_"+j + ", A_"+i
                            +", "+k+"");
                        outputKey.set("C_"+i+"_"+j);
                        // values needed to identify : matrix,
                        // position, value to multiply
                        // i.e (A,x,valA) will have to match (B,x,
                        // valB) to go valA*valB
                        outputValue.set("A,"+ k +", "+ A[i][k]);
                        context.write(outputKey, outputValue);
                    }
                }
            }
            // B
            if (matStruct[0].equals("B")) {
                // fill in matrix
                for (int i=0;i < ?? ;i++) {

                    ???
                    ???
                    ???

                }
            }
            // emit (key,value)
            for (int i=0;i< ??? ;i++) {
                for (int j=0;j< ??? ;j++) {
                    for (int k=0;k< ??? ;k++) {

                        ???
                        ???
                        ???

                    }
                }
            }
        }
    }

    public static class Reduce extends Reducer<Text, Text, Text, Text> {

        /**
         * reduce
         * @param key the key as Text
         * @param values the list of values associated to key
         * @param context the place to store outputs (goes to results)
         */
        public void reduce(Text key, Iterable<Text> values, Context context)
            throws IOException, InterruptedException {

```

```

String [] value;
int M = Integer.parseInt(context.getConfiguration().get("M"));
Float [] valsA = new Float [M];
Float [] valsB = new Float [M];
System.out.print("k="+key.toString()+"->[");
// e.g key: C_2_3 -> [("A,2,3.0");("A,2,4.0");("B,2,23.0") ...]
for (Text val : values) {
    value = val.toString().split(",");
    int loc = Integer.parseInt(value[1]);
    if (value[0].equals("A")) {
        System.out.print("A_"+value[1]+","+value[2].toString
            ());
        valsA[loc] = Float.parseFloat(value[2]);
    } else {
        System.out.print("B_"+value[1]+","+value[2].toString
            ());
        valsB[loc] = Float.parseFloat(value[2]);
    }
}
System.out.println("]");
float result = 0.0f;
float a_ij;
float b_jk;
for (int k = 0; k < ??? ; k++) {
    result += valsA[k] * valsB[k];
}
context.write(null, new Text(key.toString() + "," + Float.
    toString(result)));
}
}

public static void main(String[] args) throws Exception {

    if (args.length < 2) {
        System.err.println("Usage : hadoop jar Mm.jar Mm matrices.in res.
            out");
        System.exit(0);
    }

    Configuration conf = new Configuration();
    // A is an L-rows, M-cols matrix
    // B is an M-rows, N-cols matrix.
    conf.set("L", "2");
    conf.set("M", "5");
    conf.set("N", "3");

    Job job = new Job(conf, "MatrixMultiplication");
    job.setJarByClass(Mm.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(Text.class);

    job.setMapperClass(Map.class);
    job.setReducerClass(Reduce.class);

    job.setInputFormatClass(TextInputFormat.class);
    job.setOutputFormatClass(TextOutputFormat.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    job.waitForCompletion(true);
}

```